

MENTOR: Fixing Introductory Programming Assignments With Formula-Based Fault Localization and LLM-Driven Program Repair

Pedro Orvalho¹

¹Instituto de Investigación en Inteligencia Artificial (IIIA), Consejo Superior de Investigaciones Científicas (CSIC), Barcelona, Spain

Work done as a doctoral researcher at:
INESC-ID, Instituto Superior Técnico, Universidade de Lisboa, Portugal
CIIRC, Czech Technical University in Prague, Czechia

INESC-ID Best PhD Student 2025 Award Ceremony
Lisbon, 19 December 2025

Motivation

- The increasing demand for programming education has given rise to all kinds of online evaluations focused on introductory programming assignments (IPAs):

Motivation

- The increasing demand for programming education has given rise to all kinds of online evaluations focused on introductory programming assignments (IPAs):
 - MIT's MOOC, Introduction to CS, **reached 1.2 M enrollments** in 2018;

Motivation

- The increasing demand for programming education has given rise to all kinds of online evaluations focused on introductory programming assignments (IPAs):
 - MIT's MOOC, Introduction to CS, **reached 1.2 M enrollments** in 2018;
 - In 2020, Stanford's CS MOOC had **more than 10K students**.

Motivation

- In these courses it is a challenge to **provide personalized feedback to students.**

Motivation

- In these courses it is a challenge to **provide personalized feedback to students**.
- Providing feedback in IPAs **requires substantial time and effort by faculty**.

Automated Program Repair (APR)

Given a buggy program P_o and a set of input-output examples T (test suite).

Automated Program Repair (APR)

Given a buggy program P_o and a set of input-output examples T (test suite).


The goal of *Automated Program Repair* is to find a program P_f by **semantically change a subset S_1 of P_o 's statements** ($S_1 \subseteq P_o$) for another set of statements S_2 , s.t.,

$$P_f = ((P_o \setminus S_1) \cup S_2)$$

and


$$\forall (t_{in}^i, t_{out}^i) \in T : P_f(t_{in}^i) = t_{out}^i$$

MENTOR: Automated Feedback for Introductory Programming Exercises




- P. Orvalho, M. Janota, and V. Manquinho. MENTOR: Fixing Introductory Programming Assignments with Formula-Based Fault Localization and LLM-Driven Program Repair. In **JSS 2026**.

MENTOR




- P. Orvalho, M. Janota, and V. Manquinho. InvAASTCluster: On Applying Invariant-Based Program Clustering to Introductory Programming Assignments. In **JSS 2025**.

MENTOR



- P. Orvalho, M. Janota, and V. Manquinho. InvAASTCluster: On Applying Invariant-Based Program Clustering to Introductory Programming Assignments. In **JSS 2025**.

MENTOR



Variable Mapping

- P. Orvalho, J. Piepenbrock, M. Janota, and V. Manquinho. Graph Neural Networks For Mapping Variables Between Programs. In **ECAI 2023**.
- P. Orvalho, M. Janota, and V. Manquinho. MultiIPAs: Applying Program Transformations to Introductory Programming Assignments for Data Augmentation. In **ESEC/FSE 2022**.

Variable Mapping - Motivation

1: Function that finds and returns the maximum number among n1, n2 and n3.

```
1 int max(int n1, int n2, int n3)
2 {
3     int m = n1 > n2 ? n1 : n2;
4     return n3 > m ? n3 : m;
5 }
```

2: Function that finds and returns the maximum number among x, y and z.

```
1 int max(int x, int y, int z){
2     int m = 0;
3     m = x > m ? x : m;
4     m = y > m ? y : m;
5     return z > m ? z : m;
6 }
```

Variable Mapping: {m : m; n1 : x; n2 : y; n3 : z}.

Program Representation


3: An expression that uses int variables a and b, previously declared in the program.

```
1  {  
2      // a and b are ints  
3      a = a - b;  
4 }
```

Types of edges:


AST	
Read	
Write	
Sibling	
Chronological	

Variable Node



(a) Our program representation.

MENTOR



Fault Localization

- P. Orvalho, M. Janota, and V. Manquinho. CFAULTS: Model-Based Diagnosis for Fault Localization in C with Multiple Test Cases. In **FM 2024**.


Fault Localization

- Given a buggy program, *fault localization (FL)* involves identifying locations in the program that could cause a faulty behaviour (bug).




Formula-Based Fault Localization (FBFL)

- FBFL methods encode the localization problem into **several optimization problems** to identify a minimal set of bugs (diagnoses).



MENTOR



Program Repair

- P. Orvalho, M. Janota, and V. Manquinho. Counterexample Guided Program Repair Using Zero-Shot Learning and MaxSAT-based Fault Localization. In **AAAI 2025**.

Motivation

4: Semantically incorrect program. Faults: {4,8}.

```
1  int main(){ //finds max of 3 nums
2      int f,s,t;
3      scanf("%d%d%d",&f,&s,&t);
4      if (f < s && f >= t)
5          printf("%d",f);
6      else if (s > f && s >= t)
7          printf("%d",s);
8      else if (t < f && t < s)
9          printf("%d",t);
10
11     return 0;
12 }
```

Motivation

5: Semantically incorrect program. Faults: {4,8}.

```
1  int main(){ //finds max of 3 nums
2      int f,s,t;
3      scanf("%d%d%d",&f,&s,&t);
4      if (f < s && f >= t)
5          printf("%d",f);
6      else if (s > f && s >= t)
7          printf("%d",s);
8      else if (t < f && t < s)
9          printf("%d",t);
10
11     return 0;
12 }
```

Motivation

6: Semantically incorrect program. Faults: {4,8}.

```
1  int main(){ //finds max of 3 nums
2      int f,s,t;
3      scanf("%d%d%d",&f,&s,&t);
4      if (f < s && f >= t)
5          printf("%d",f);
6      else if (s > f && s >= t)
7          printf("%d",s);
8      else if (t < f && t < s)
9          printf("%d",t);
10
11      return 0;
12 }
```

LLMs for code (LLMCs)

- GRANITE and CODEGEMMA **cannot** fix the buggy program within 90 secs;

Motivation

7: Semantically incorrect program. Faults: {4,8}.

```
1  int main(){ //finds max of 3 nums
2      int f,s,t;
3      scanf("%d%d%d",&f,&s,&t);
4      if (f < s && f >= t)
5          printf("%d",f);
6      else if (s > f && s >= t)
7          printf("%d",s);
8      else if (t < f && t < s)
9          printf("%d",t);
10
11      return 0;
12 }
```

LLMs for code (LLMCs)

- GRANITE and CODEGEMMA **cannot** fix the buggy program within 90 secs;
- Even if we provide the assignment's **description and IO tests**.

Program Sketches


8: Semantically incorrect program. Faults :{4,8}.

```
1  int main(){ //finds max of 3 nums
2      int f,s,t;
3      scanf("%d%d%d",&f,&s,&t);
4      if (f < s && f >= t)
5          printf("%d",f);
6      else if (s > f && s >= t)
7          printf("%d",s);
8      else if (t < f && t < s)
9          printf("%d",t);
10
11     return 0;
12 }
```

9: Program sketch with holes.

```
1  int main(){
2      int f,s,t;
3      scanf("%d%d%d",&f,&s,&t);
4      @ HOLE 1 @
5          printf("%d",f);
6      else if (s > f && s >= t)
7          printf("%d",s);
8      @ HOLE 2 @
9          printf("%d",t);
10
11     return 0;
12 }
```

Counterexample Guided Automated Repair




LLM-Driven APR with CFaults

LLMs	Prompt Configurations			
	De-TS	De-TS-CE	Sk_De-TS	Sk_De-TS-CE
CodeGemma	41.7%	42.3%	47.7%	48.1%
CodeLlama	34.4%	34.9%	40.0%	39.2%
Gemma	34.7%	34.4%	37.2%	37.3%
Granite	43.7%	43.6%	48.3%	47.6%
Llama3	39.4%	41.2%	40.4%	41.3%
Phi3	34.5%	34.2%	38.2%	37.4%
Verifix	6.3%			
Clara	34.6%			

Table 1: The number of programs fixed by each LLM under various configurations. Mapping abbreviations to configuration names: **De** - IPA *Description*, **TS** - *Test Suite*, **CE** - *Counterexample*, **SK** - *Sketches*.

MENTOR: Automated Feedback for Introductory Programming Exercises



- P. Orvalho, M. Janota, and V. Manquinho. MENTOR: Fixing Introductory Programming Assignments with Formula-Based Fault Localization and LLM-Driven Program Repair. In **JSS 2026**.

Takeaway Message

- We tackle the APR problem using an **LLM-Driven Counterexample Guided Inductive Synthesis (CEGIS) approach** [Solar-Lezama et al., 2006];

Takeaway Message

- We tackle the APR problem using an **LLM-Driven Counterexample Guided Inductive Synthesis (CEGIS) approach** [Solar-Lezama et al., 2006];
- We employ **Logic-based Fault Localization to guide and minimize LLMs' patches** to incorrect programs by feeding them bug-free program sketches;

Takeaway Message

- We tackle the APR problem using an **LLM-Driven Counterexample Guided Inductive Synthesis (CEGIS) approach** [Solar-Lezama et al., 2006];
- We employ **Logic-based Fault Localization to guide and minimize LLMs' patches** to incorrect programs by feeding them bug-free program sketches;
- With our approach **all six evaluated LLMs fix more programs and produce smaller patches** than other configurations and symbolic tools;

Takeaway Message

- We tackle the APR problem using an **LLM-Driven Counterexample Guided Inductive Synthesis (CEGIS) approach** [Solar-Lezama et al., 2006];
- We employ **Logic-based Fault Localization to guide and minimize LLMs' patches** to incorrect programs by feeding them bug-free program sketches;
- With our approach **all six evaluated LLMs fix more programs and produce smaller patches** than other configurations and symbolic tools;
- All our code is available on GitHub and on Zenodo.

Obrigado!
Thank you!

Obrigado!
Thank you!

My sincere thanks to INESC-ID,

Obrigado!
Thank you!

My sincere thanks to INESC-ID,
and to my advisors, Vasco and Mikoláš!

Pedro Orvalho

Thank you!



<https://pmorvalho.github.io>

References



Reiter, Raymond (1987)

A Theory of Diagnosis from First Principles.

Artif. Intell. 1987.



Do, Hyunsook and Elbaum, Sebastian G. and Rothermel, Gregg (2005)

Supporting Controlled Experimentation with Testing Techniques: An Infrastructure and its Potential Impact.

Empir. Softw. Eng. 2005.



Jose, Manu and Majumdar, Rupak (2011)

Cause clue clauses: error localization using maximum satisfiability.

PLDI 2011.



Armando Solar-Lezama and Liviu Tancau and Rastislav Bodík and Sanjit A. Seshia and Vijay A. Saraswat (2018)

Combinatorial sketching for finite programs.

ASPLOS 2006.

References

-  Ignatiev, Alexey and Morgado, António and Weissenbacher, Georg and Marques-Silva, João (2019) Model-Based Diagnosis with Multiple Observations.
IJCAI 2019.
-  Orvalho, Pedro and Janota, Mikolas and Manquinho, Vasco (2024) C-Pack of IPAs: A C90 Program Benchmark of Introductory Programming Assignments.
Automated Program Repair (APR) 2024.
-  Lamraoui, Si-Mohamed and Nakajima, Shin (2016) A Formula-based Approach for Automatic Fault Localization of Multi-fault Programs.
J. Inf. Process. 24(1), 88 – 98.
-  The Guardian UK - Crowdstrike Meltdown
<https://www.theguardian.com/technology/article/2024/jul/24/crowdstrike-outage-companies-cost>.
The Guardian UK.
-  Ahmed, Umair Z and Fan, Zhiyu and Yi, Jooyong and Al-Bataineh, Omar I and Roychoudhury, Abhik (2022) Verifix: Verified repair of programming assignments.
TOSEM 22 12(3), 45 – 678.

References



Orvalho, Pedro and Janota, Mikoláš and Manquinho, Vasco (2022)

MulIPAs: Applying Program Transformations to Introductory Programming Assignments for Data Augmentation.
ESEC/FSE 2022.



Gulwani, Sumit and Radiček, Ivan and Zuleger, Florian (2018)

Automated clustering and program repair for introductory programming assignments.
PLDI 18 52(4), 465 – 480.



Orvalho, Pedro and Janota, Mikolas and Manquinho, Vasco (2024)

CFaults: Model-Based Diagnosis for Fault Localization in C with Multiple Test Cases.
Formal Methods (FM) 2024.



Orvalho, Pedro and Janota, Mikolas and Manquinho, Vasco (2025)

Counterexample Guided Program Repair Using Zero-Shot Learning and MaxSAT-based Fault Localization.
AAAI 2025.